

# Making the Machine Learning Reconstruction Pipeline Scalable

Jack Carlton  
University of Kentucky

# Breaking down “Scalability”

I choose to break down scalability into three categories:

## 1. Data

- a. ML pipeline does not change behavior as data set grows
  - i. Still scales in execution time

## 2. Compute

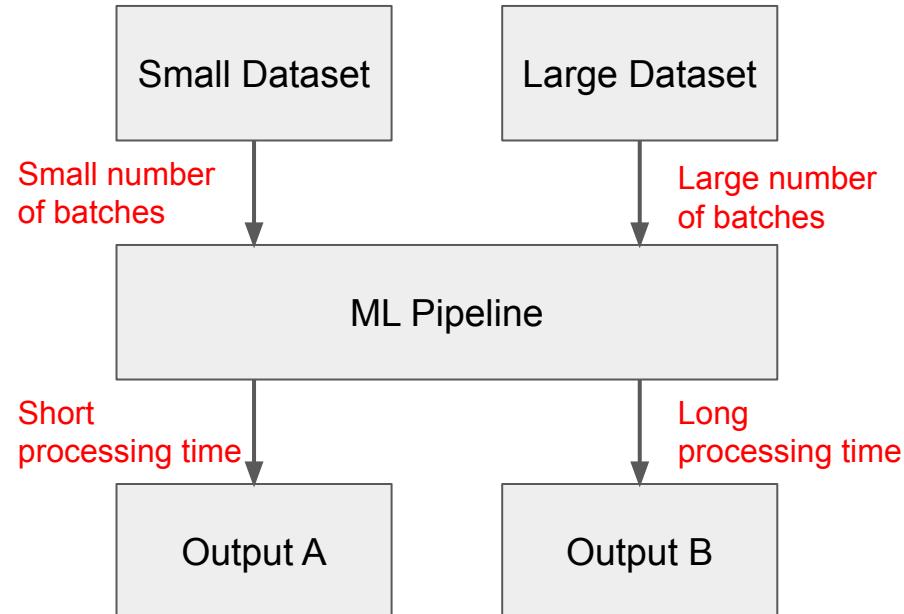
- a. ML pipeline does not change behavior as compute power is changed
  - i. Still scales in execution time
  - ii. Ex. Pipeline runs on developer's laptop or a CPU/GPU cluster

## 3. Codebase

- a. ML pipeline does not (greatly) change behavior as complexity grows
  - i. New models, stages, or data products require additional code, not (major) code rewrites
  - ii. Iteration speed does not (greatly) degrade with system size (i.e. keep things modular!)

# What Do We Mean by “Scaling” (Data)

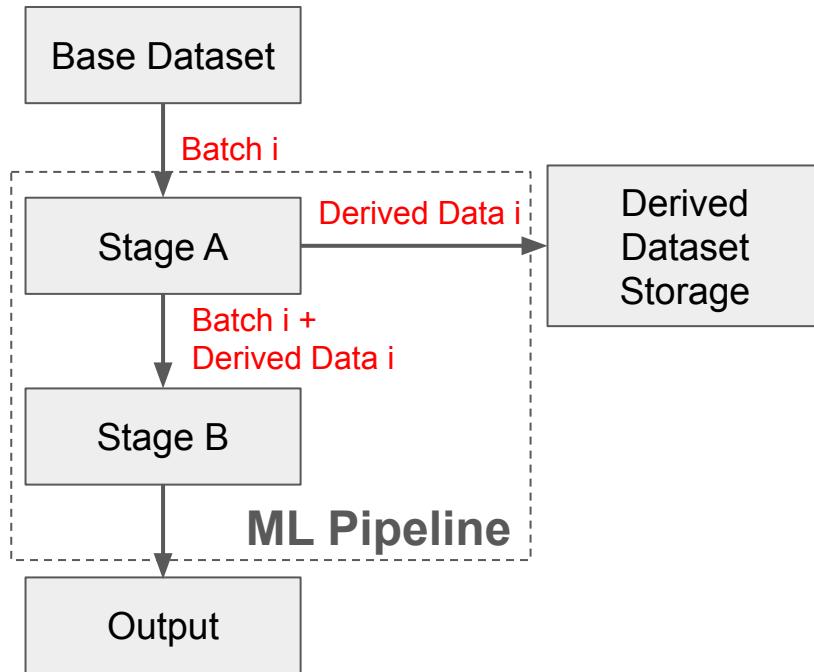
- **Data volume should be able to grow without changing how the pipeline behaves**
- Adding the following should not cause pipeline behavior changes:
  - More events
  - More derived data products
    - predictions, masks, regressions, etc.
  - More passes over the same data
  - Larger event representations
- Implications:
  - Memory usage must not grow with dataset size



**Ideal Behavior for Different Sized Datasets**

# Techniques to Ensure Scalability (Data)

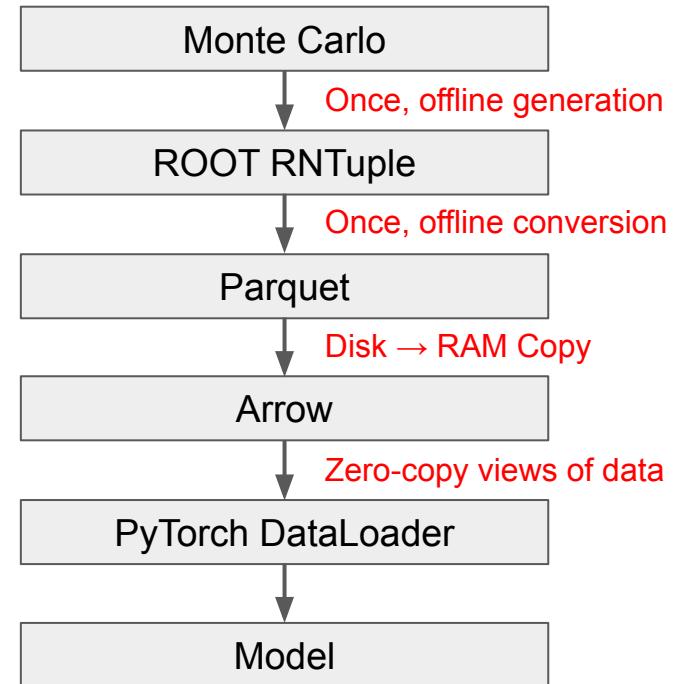
- Stream data in bounded batches
  - RAM usage should not scale with data set size
  - Allows RAM usage to be “tunable”
- Base data should be immutable
  - No modifications or extensions to ML pipeline input data files
- Separate derived data products
  - Predictions, masks, and regressions are produced as independent datasets
- Reference data by IDs, not by object
  - When passing data modules, use file paths or map IDs not in memory collections
  - Similar to why we use pointers in C++



**Simple Example Pipeline Including Derived Datasets**

# Technologies for Scalability (Data)

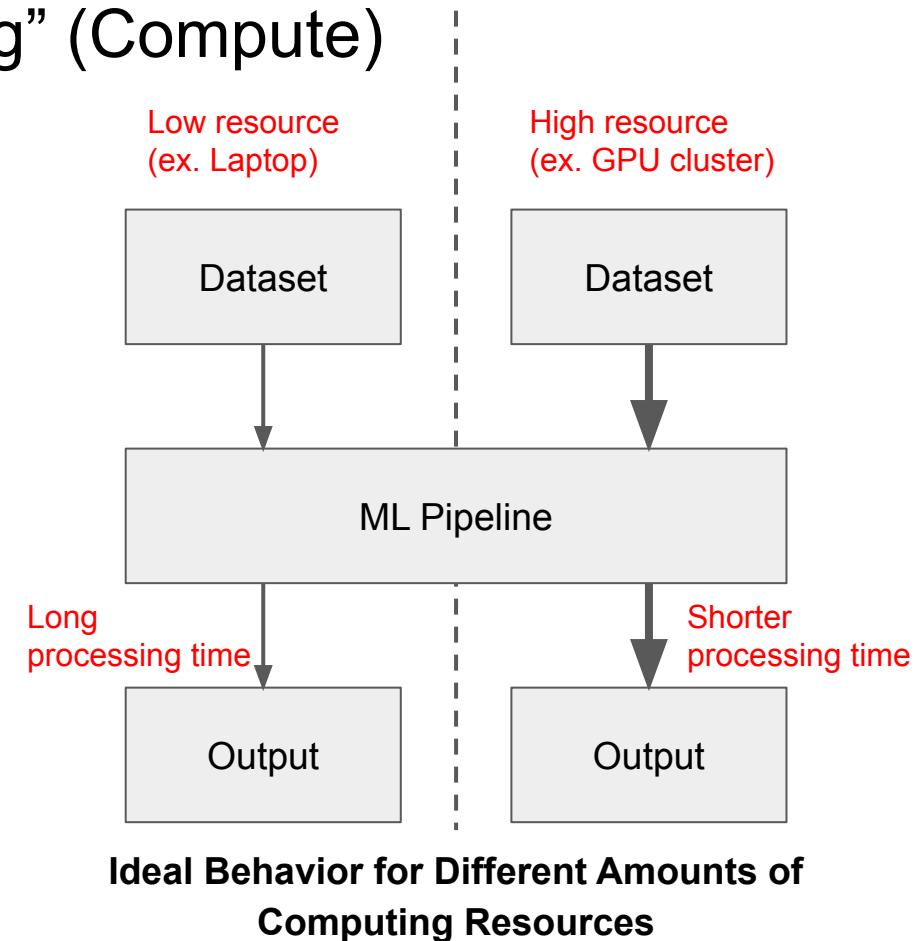
- ROOT RNTuple (or similar)
  - Holds the necessary data from the monte carlo
- Apache Parquet (columnar, on disk)
  - ML-native working datasets, append-only, shardable (aka “batchable”)
  - Can key rows by event ID, enabling batch-scoped joins for derived data products
- Apache Arrow
  - Single copy from disk into RAM, then zero-copy views all the way to Torch
- PyTorch tensors
  - Execution format for models
- Pytorch DataLoader
  - Handles batching, shuffling, parallel loading, prefetching, enforcing memory bounds



**Flow of Data From Monte Carlo to a Model in the ML pipeline**

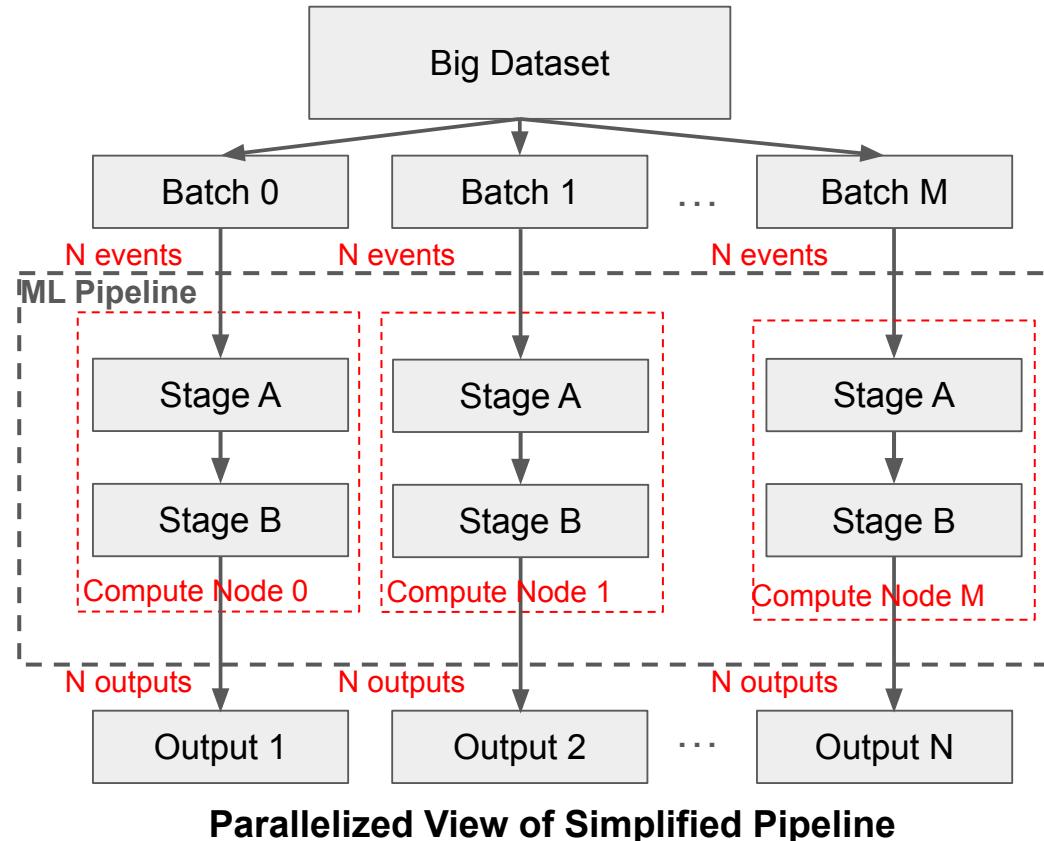
# What Do We Mean by “Scaling” (Compute)

- Computation resources should be able to grow without changing how the pipeline behaves
- Adding the following should not cause pipeline behavior changes:
  - Number of GPUs
  - Number of CPU cores/threads
  - Node count
  - Memory Capacity
- Implications:
  - Algorithms must be agnostic to resources
    - Caveat: PyTorch and other frameworks may change their behavior for different devices/device counts for efficiency



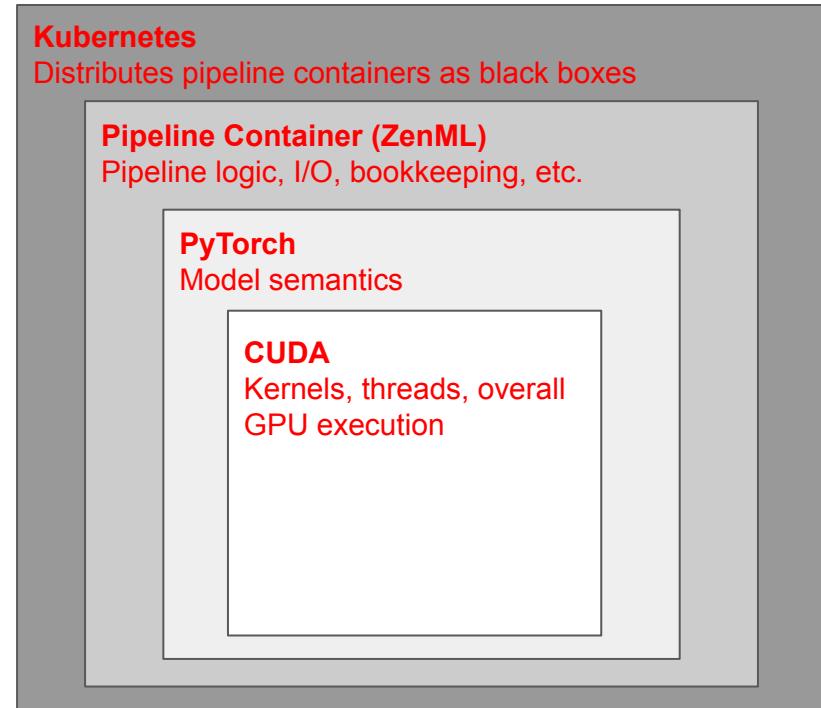
# Techniques to Ensure Scalability (Compute)

- Make algorithms resource-agnostic
  - No logic branches based on GPU count, core count, node count, or memory size, etc.
- Make algorithms easily parallelizable
  - Decompose computation into independent, composable units
  - Avoid global state dependencies in pipeline stages
  - Avoid “synchronization points”; i.e. points where models must make inferences on whole datasets



# Technologies for Scalability (Compute)

- **PyTorch**
  - Standard framework for modern ML models
  - Resource-agnostic execution model
- **CUDA**
  - Operates purely at the level of memory, kernels, and execution, independent of algorithm or pipeline semantics
  - Provides an API for launching and coordinating large numbers of parallel threads on GPUs
- **Kubernetes**
  - Schedules identical pipeline executions onto available compute nodes
  - Scales how many pipelines run concurrently, not what they do
  - Handles retry logic and resource limits

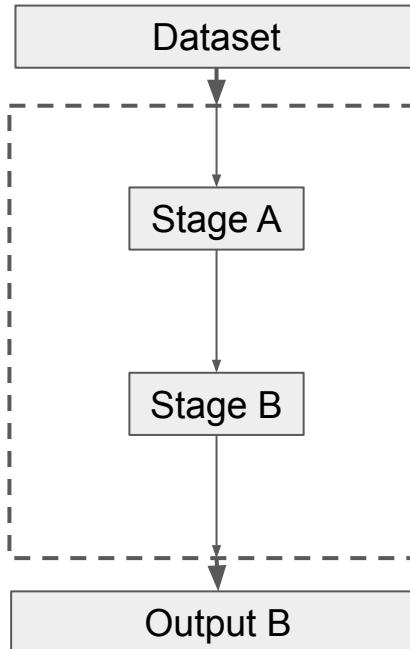


**Simplified “Scope” Of Technologies**  
**Outer Technologies Manage Inner Technologies**

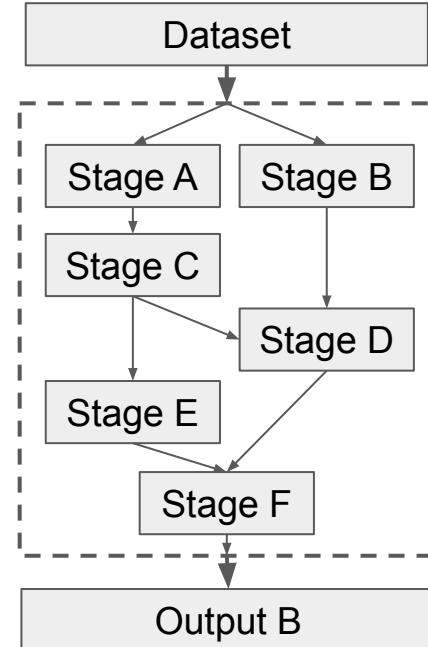
# What Do We Mean by “Scaling” (Codebase)

- **Codebase complexity should be able to grow without changing how the system behaves**
- Adding the following should not cause system wide behavioral changes:
  - More pipeline stages
  - More models / algorithms
  - More configuration options
- Implications:
  - New functionality should be added by extensions, not modification
  - System behavior should be locally understood (modularity)
  - Existing code should not require global refactoring to evolve

“Simple” pipeline with few stages



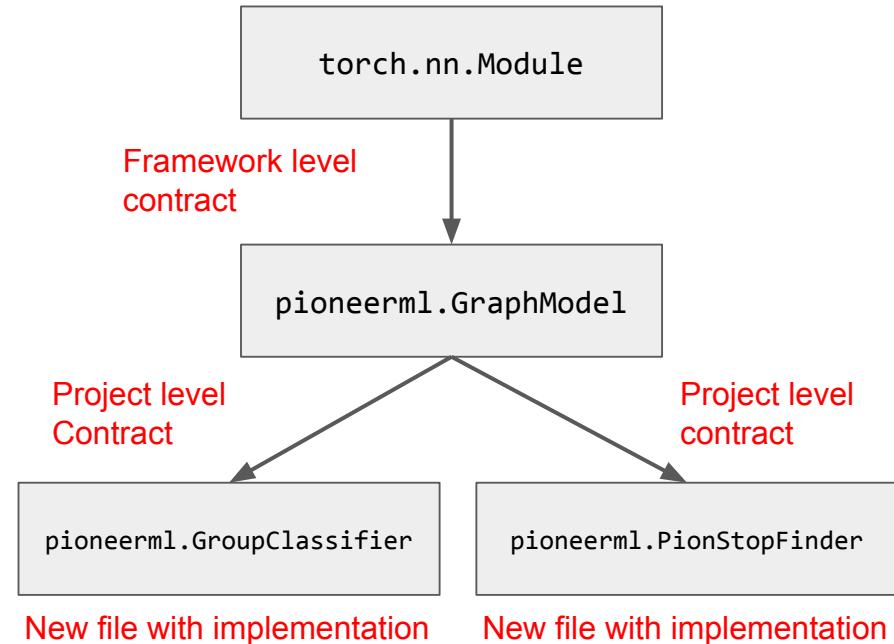
“Complex” Pipeline with many stages



The “Global” Structure Should Not Change As Complexity Increases

# Techniques to Ensure Scalability (Codebase)

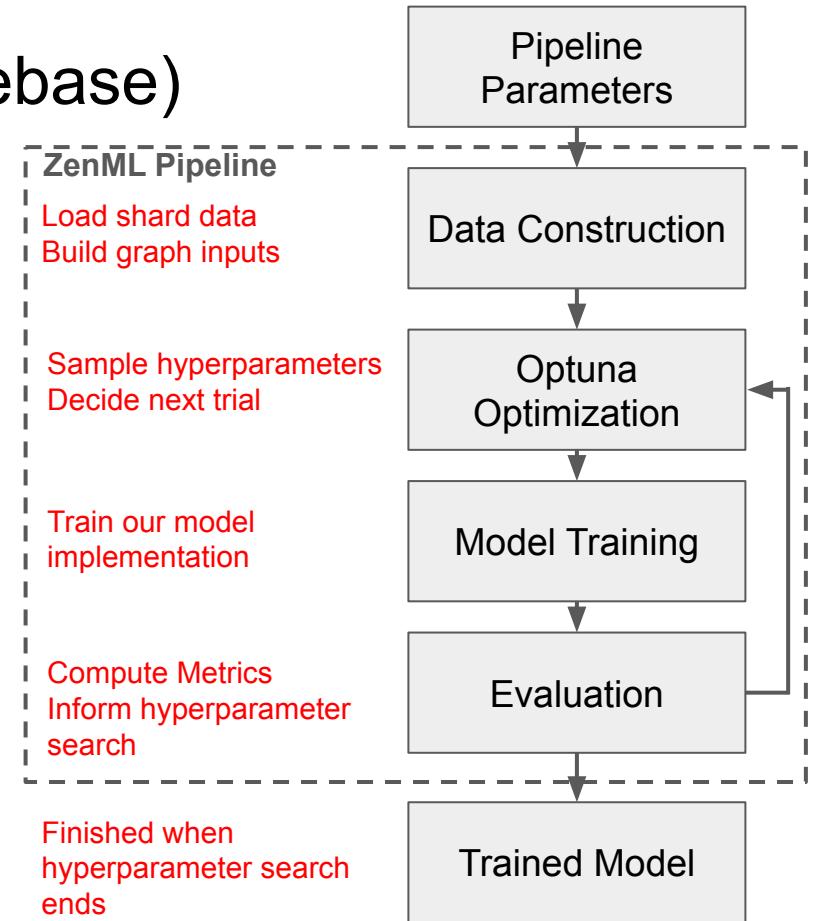
- Use abstraction where appropriate
  - Define stable base interfaces / abstract classes
  - Add new functionality by extending, not rewriting
- Avoid global coupling
  - Modules depend only on explicit inputs
  - Each module manages its own immutable local state
- Isolate responsibilities to achieve modularity
  - Each module has a single, well-defined responsibility
  - Changes remain local to the owning module



**Simple Example of Abstraction For Adding New Models**

# Technologies for Scalability (Codebase)

- PyTorch
  - Provides a standard base abstraction classes for many model types (ex. `nn.Module`)
  - Enforces a consistent model interface (ex. `forward` method)
- ZenML
  - Encodes pipelines as composable, declarative units and orchestrates execution
  - Allows pipelines to grow by adding or reordering steps; easy to add new pipelines
  - Manages pipeline state, artifacts, and execution metadata outside user code
- Optuna
  - Isolates hyperparameter search code
  - Enables experimentation without modifying core implementations



Simplified Example Pipeline for Training Models

# Auxiliary Slides

# What is Apache Parquet?

- **Apache Parquet is a columnar, on-disk data format that is widely used in ML workloads**
- What parquet does
  - Columnar storage → read only the columns (features) your model needs
    - Allows one to efficiently assign a subset of columns as inputs and another subset of columns as targets
  - Supports nested and variable-length fields
  - Data schema is embedded in the file, not inferred by code (ex. Not like numpy, where code defines `dtype` parameter)
- Why this matters for scalability
  - Efficient I/O for large datasets through compression and encoding
  - New models using different subsets of the data becomes trivial

Columns:

---

```
event_id
theta
phi
pion_stop_x
pion_stop_y
pion_stop_z
total_pion_energy
...
hit_coord→ list < float >
hit_z→ list < float >
hit_energy→ list < float >
hit_pdg_mask→ list < int >
```

Example `time_groups.parquet` file structure

# What is Batching?

- **Batching means processing a fixed-size subset of events at a time, rather than the full dataset**
- What batching does
  - Groups individual events into batches of size N
  - Each batch is processed independently
  - Batches are discarded after use
- Why this matters for scalability
  - Memory usage is bounded by batch size (for single batch process)
  - Dataset size does not affect RAM usage
  - Enables streaming over arbitrarily large datasets

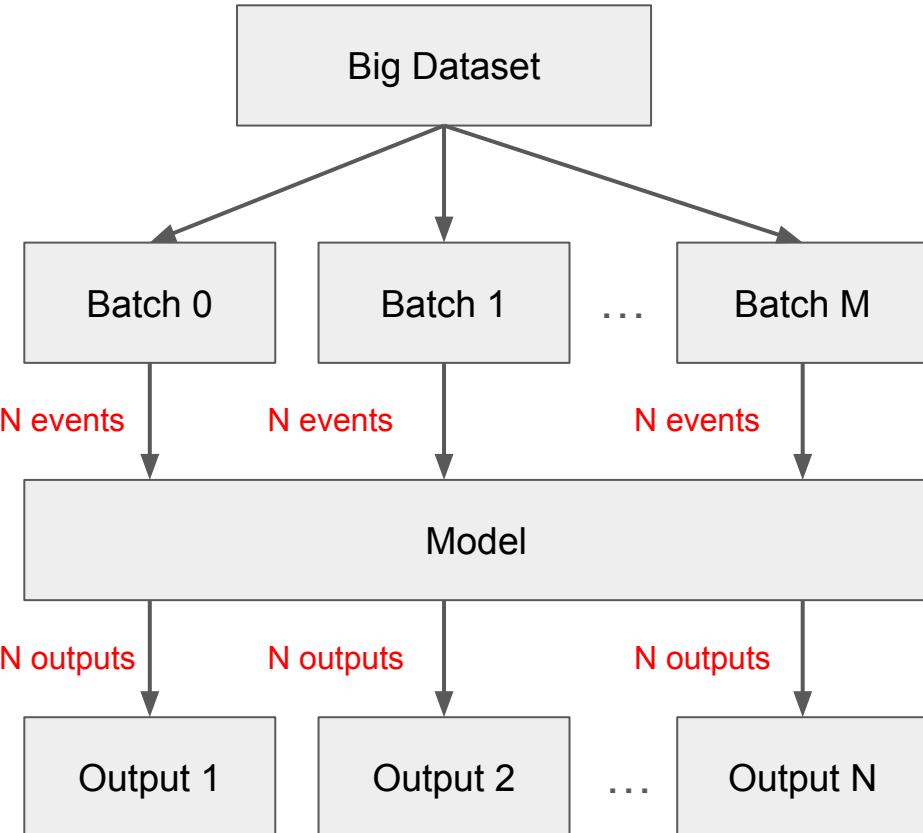
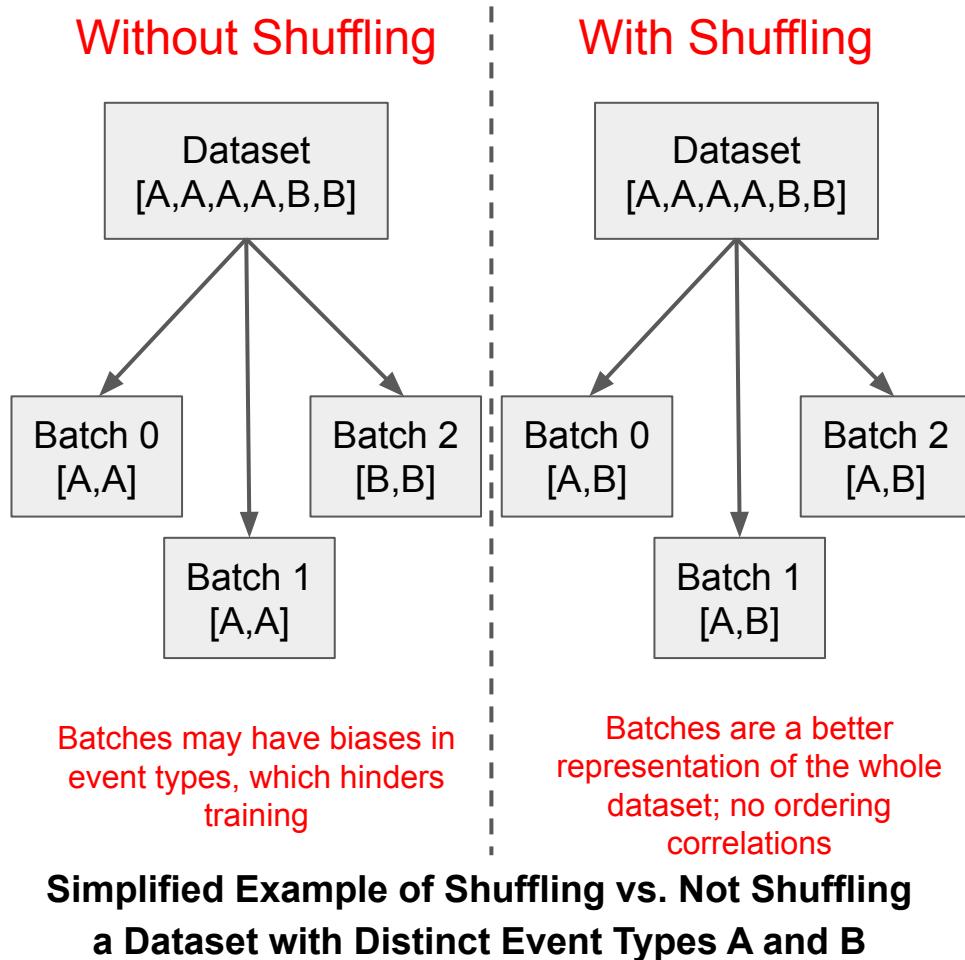


Diagram that Shows how Batching Splits Data

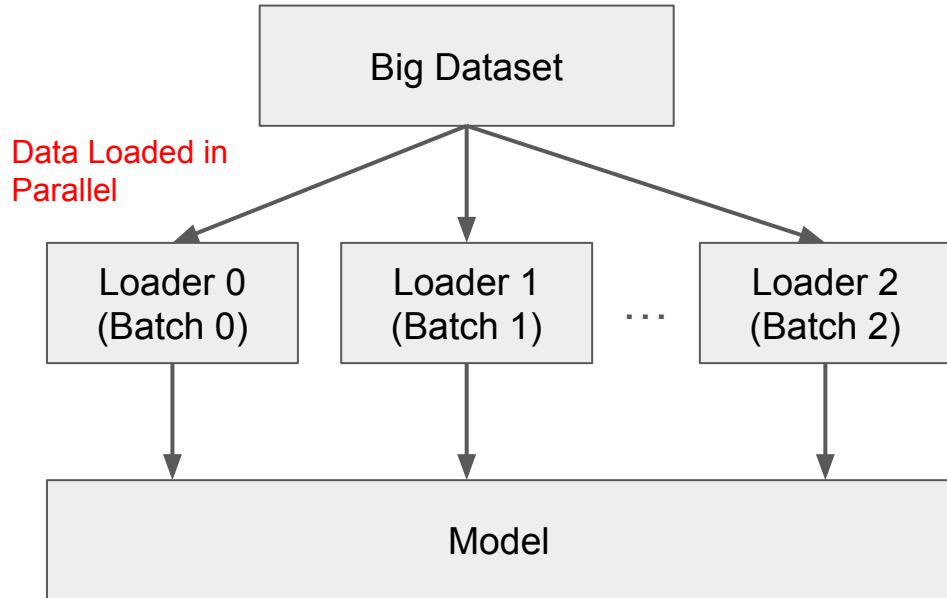
# What is Shuffling?

- **Shuffling changes the order in which events are seen, without changing the data itself**
- What shuffling does
  - Randomizes event order before forming batches
  - Ensures batches contain a mix of events
  - Changes between epochs (or passes over the data)
- Why this matters for scalability
  - Prevents bias from data ordering
  - Improves statistical independence between batches
  - Allows repeated passes over large datasets without correlation artifacts



# What is Parallel Loading?

- **Parallel loading means loading multiple batches concurrently, using multiple workers**
- What parallel loading does
  - Multiple workers read and prepare batches simultaneously
  - The model always has a batch ready to process
  - Data loading is decoupled from model execution
- Why this matters for scalability
  - Helps prevent the model from waiting on disk I/O
  - Improves hardware utilization (especially GPUs)



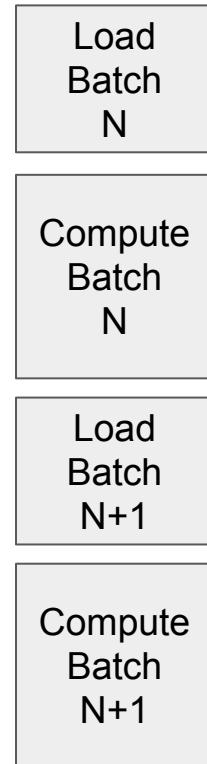
Note: Parallel loading does **not** necessarily mean parallel model execution

Parallel Data Loading Diagram

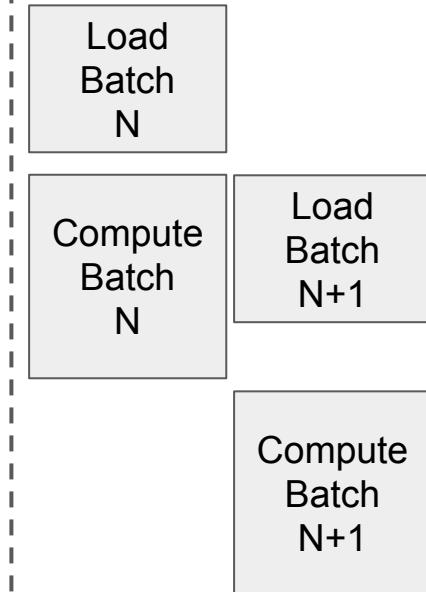
# What is Prefetching?

- **Prefetching means loading future batches while the current batch is being processed**
- What prefetching does
  - While the model computes on batch N the next batch (N+1) is loaded in the background
  - When computation finishes, the next batch is already ready
- Why this matters for scalability
  - Helps prevent the model from waiting on disk I/O

## Without Prefetching



## With Prefetching



Subsequent batches loaded in parallel with compute of previous batch. Effectively utilizing different hardwares

Simplified Prefetching Example Diagram

# What is Enforcing Memory Bounds?

- **Enforcing memory bounds means placing a hard limit on how much data can be in memory at once**
- What enforcing memory bounds does
  - Maximum in-flight data size is fixed
  - Batch creation is throttled when memory is full
  - Makes memory usage predictable and stable
- Why this matters for scalability
  - Dataset size does not affect RAM usage
  - Enables streaming over arbitrarily large datasets

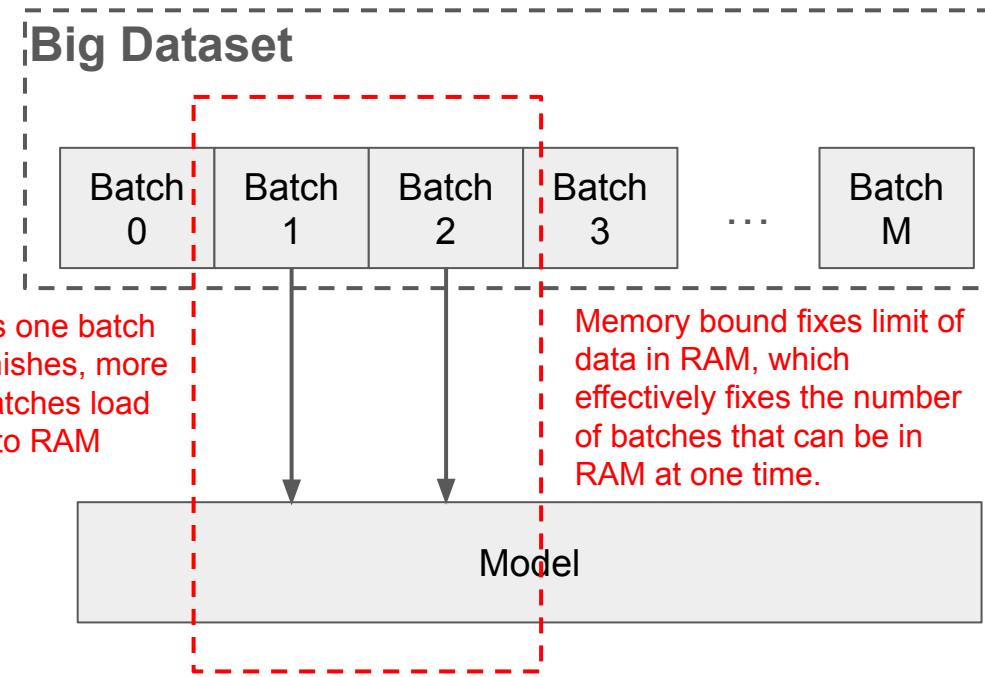
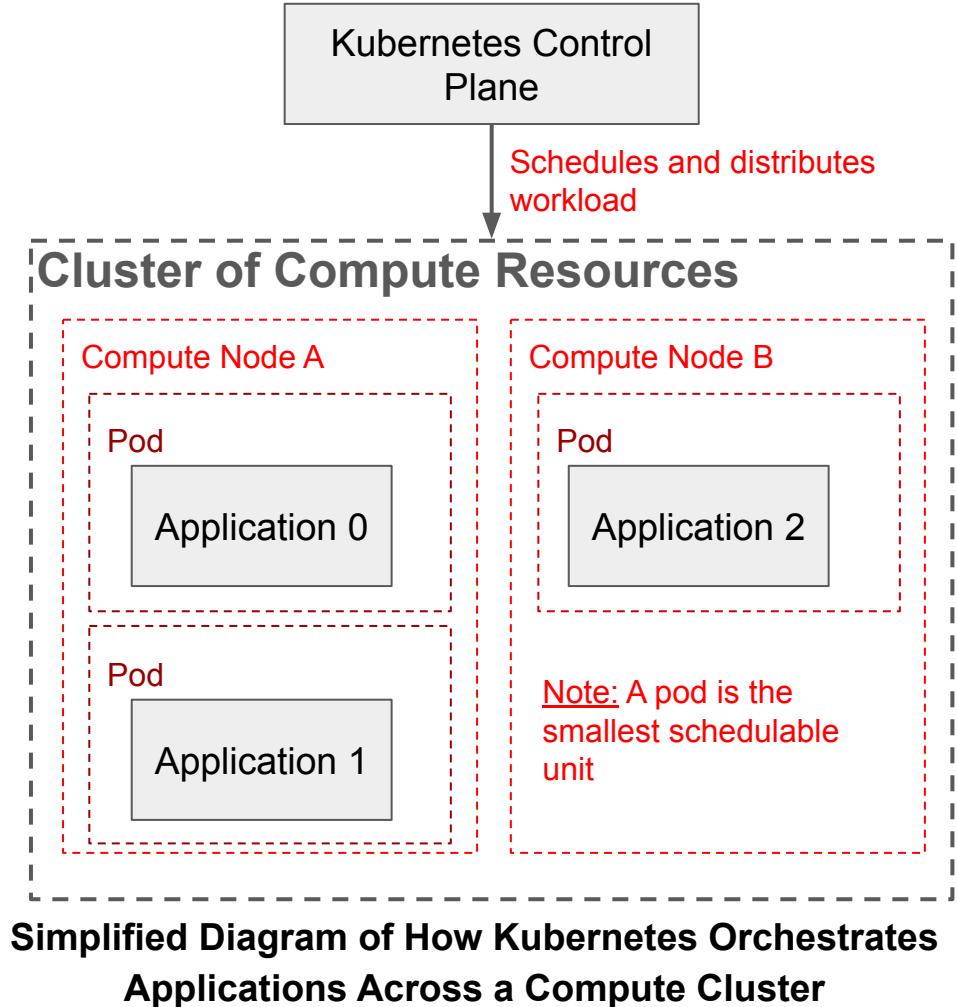


Diagram Showing How Memory Bounds Enforce a Limited Number of Batches Loaded in RAM

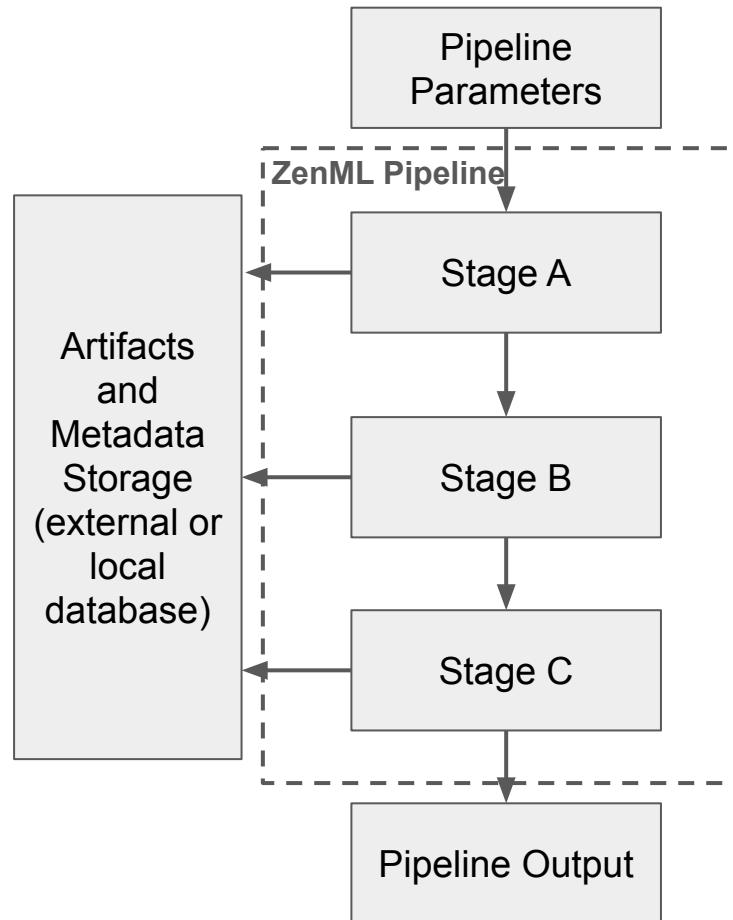
# What is Kubernetes (K8s)?

- **A system for running and managing containerized workloads across shared compute resources**
- What Kubernetes does
  - Schedules containers onto available compute nodes
  - Enforces resource limits (CPU, GPU, memory) per container
  - Isolates workloads from one another
  - Handles restarts and retries on failure
- Why this matters for scalability
  - Enables concurrent execution of many independent workloads
  - Prevents resource contention between workloads



# What is ZenML?

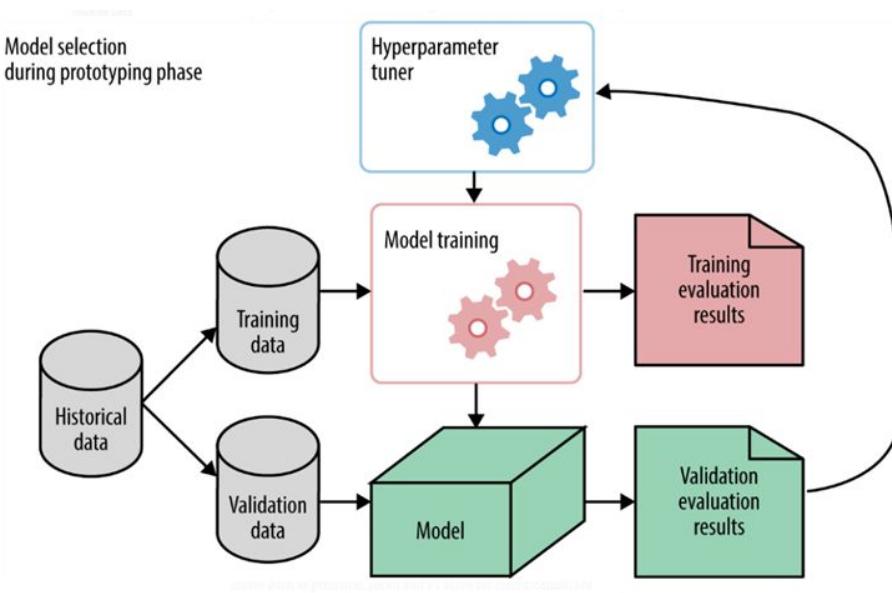
- **A framework for defining, orchestrating, and executing machine-learning pipelines with explicit steps, artifacts, and metadata**
- What ZenML does
  - Encodes workflows as composable, declarative pipelines
  - Orchestrates execution order and dependencies
  - Integrates with execution backends (local, containers, clusters) without changing user code
- Why this matters for scalability
  - Pipelines grow by adding or rearranging stages, not rewriting logic
  - Enables reproducibility, versioning, and parallel development
  - State and artifacts are managed outside user code



Example ZenML pipeline diagram

# What is Optuna?

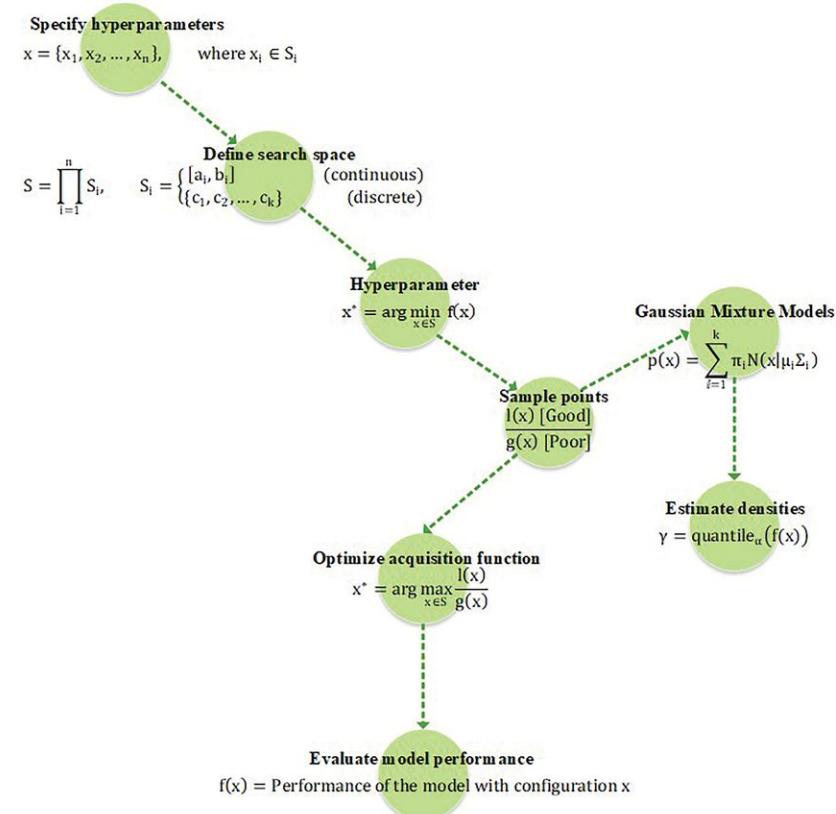
- Many hyper parameters means manual tuning is bad
  - Too slow
  - Suboptimal tuning means you spend more resources (time, computing power, etc.) training unused models
- Optuna is a python package that solves this problem
  - Framework for optimization black box objective functions
  - Technically not an ML package, but rather a package that supports many optimized sampling strategies



Example Optuna workflow diagram

# How Optuna Works

- Optuna supports many sampling algorithms, examples:
  - Grid search
  - Random search
  - Gaussian process-based Bayesian optimization
- For single object functions, the default for Optuna is Tree-Structured Parzen Estimator (TPE)



TPE flow diagram (in a nutshell)

# What is a Tree-Structured Parzen Estimator (TPE) (Part I)

- Given a (possibly stochastic) black box function you want to minimize (ex. model loss)

$$y = f(\theta)$$

- Hyperparams  $\equiv \theta$
- Define the following:

$\gamma \equiv$  "good"/"bad" fraction (say, 0.2)

$y^*$  is chosen such that  $P(y < y^*) = \gamma$   
so  $y^*$  is in, say, the best 20% of losses

$\max(y^* - y, 0) \equiv$  improvement

- Then, for any given theta we can define *expected improvement*

$$\text{EI}(\theta) \equiv \mathbb{E}[\max(y^* - y, 0) \mid \theta] = \int_{-\infty}^{y^*} (y^* - y) p(y \mid \theta) dy$$

- Goal: Maximize expected improvement

# What is a Tree-Structured Parzen Estimator (TPE) (Part II)

- For a 1 dimensional  $y$ , it turns out to be easier to work with  $p(\theta | y)$ , we can invert using Bayes' rule:

$$p(y | \theta) = \frac{p(\theta | y) p(y)}{p(\theta)}$$

- And substitute

$$\text{EI}(\theta) \propto \int_{-\infty}^{y^*} (y^* - y) \frac{p(\theta | y)}{p(\theta)} p(y) dy$$

- Since we don't know every value of  $y$  this becomes an impossible task. We must make an approximation by dividing into "good" and "bad" distributions

$$p(\theta | y) \approx \begin{cases} p(\theta | y < y^*) \equiv \ell(\theta), & y < y^* \text{ (good)} \\ p(\theta | y \geq y^*) \equiv g(\theta), & y \geq y^* \text{ (bad)} \end{cases}$$

$$p(\theta) \approx \ell(\theta) \int_{y < y^*} p(y) dy + g(\theta) \int_{y \geq y^*} p(y) dy = \gamma l(\theta) + (1 - \gamma) g(\theta)$$

# What is a Tree-Structured Parzen Estimator (TPE) (Part III)

- The integral, by construction, only cares about the “good” region, so the integral simplifies to

$$\text{EI}(\theta) \propto \frac{\ell(\theta)}{p(\theta)} \int_{-\infty}^{y^*} (y^* - y) p(y) dy$$

- But the integral is now constant in theta! So we have:

$$\text{EI}(\theta) \propto \frac{\ell(\theta)}{p(\theta)}$$

- Where  $p(\theta) = \gamma \ell(\theta) + (1 - \gamma) g(\theta)$  so we can write:

$$\text{EI}(\theta) \propto \frac{\ell(\theta)}{\gamma \ell(\theta) + (1 - \gamma) g(\theta)}$$

# What is a Tree-Structured Parzen Estimator (TPE) (Part IV)

- But  $\gamma$  is fixed, so

$$\begin{aligned}\arg \max_{\theta} EI(\theta) &= \arg \max_{\theta} \frac{\ell(\theta)}{\gamma \ell(\theta) + (1 - \gamma) g(\theta)} \\ &= \arg \max_{\theta} \frac{\ell(\theta)/g(\theta)}{\gamma \ell(\theta)/g(\theta) + (1 - \gamma)} \\ &= \arg \max_{\theta} \frac{\ell(\theta)}{g(\theta)}\end{aligned}$$

- Where the final step is because  $x/(bx+c)$  is monotonic in  $x$  for  $x > 0$ , let  $x = l/g$
- In other words, we just need to find  $\arg \max_{\theta} \frac{\ell(\theta)}{g(\theta)}$  which is doable via algorithm!

# What is a Tree-Structured Parzen Estimator (TPE) (Part V)

- Now to define the algorithm, first we observe T ( $\sim 10$ ) trials randomly:

**Observed trials:**  $\mathcal{D} = \{(\theta^{(i)}, y^{(i)})\}_{i=1}^T$

**Quantile threshold:**  $P(y < y^*) = \gamma$

$$\mathcal{D}_{\text{good}} = \{\theta^{(i)} : y^{(i)} < y^*\}$$

$$\mathcal{D}_{\text{bad}} = \{\theta^{(i)} : y^{(i)} \geq y^*\}$$

- From this data, we want to build:

$$\ell(\theta) \equiv p(\theta \mid y < y^*)$$

$$g(\theta) \equiv p(\theta \mid y \geq y^*)$$

- So we define

$$\theta \equiv (\theta_1, \theta_2, \dots, \theta_d)$$

- And assume:

**Independence assumption:**  $p(\theta \mid y < y^*) \approx \prod_{k=1}^d p(\theta_k \mid y < y^*)$

# What is a Tree-Structured Parzen Estimator (TPE) (Part VI)

- This allows us to write:

$$\ell(\theta) \approx \prod_{k=1}^d \hat{p}_{\text{good}}(\theta_k)$$

$$g(\theta) \approx \prod_{k=1}^d \hat{p}_{\text{bad}}(\theta_k)$$

- We can fit to our samples to get a continuous distribution spaces for each param

$$\hat{p}_{\text{good}}(\theta_k) \leftarrow \text{fit to } \{\theta_k^{(i)} : \theta^{(i)} \in \mathcal{D}_{\text{good}}\}$$

$$\hat{p}_{\text{bad}}(\theta_k) \leftarrow \text{fit to } \{\theta_k^{(i)} : \theta^{(i)} \in \mathcal{D}_{\text{bad}}\}$$

- Then we sample from the “good” spaces for M candidates index by m

$$\tilde{\theta}_{m,k} \sim \hat{p}_{\text{good}}(\theta_k), \quad k = 1, \dots, d$$

- And finally, choose our next theta, add it to the data set, and repeat

$$\theta^{(t)} = \arg \max_{\tilde{\theta}_m} \frac{\ell(\tilde{\theta}_m)}{g(\tilde{\theta}_m)} = \prod_{k=1}^d \frac{\hat{p}_{\text{good}}(\tilde{\theta}_{m,k})}{\hat{p}_{\text{bad}}(\tilde{\theta}_{m,k})}$$

# What is a Tree-Structured Parzen Estimator (TPE) (Part VII)

- How do we obtain our fits from the data?
- Uses [Kernel Density Estimation](#) (KDE)
  - The actual fit is done when determining each “gaussian kernel” (or sigma)
  - Likelihood-optimal is expensive  $\sim O(n^2d)$ 
    - $n = \# \text{ samples}$
    - $d = \# \text{ hyperparameters}$
  - Optuna uses the “heuristic” version, which requires choosing some value for  $c$ .

Given samples:  $\{\theta_k^{(i)}\}_{i=1}^n$ ,  $\theta_k^{(i)} \in \mathbb{R}$

**Kernel density estimate:**

$$\hat{p}(\theta_k) = \frac{1}{n} \sum_{i=1}^n \mathcal{K}\left(\frac{\theta_k - \theta_k^{(i)}}{h_k}\right)$$

**Gaussian kernel:**

$$\mathcal{K}(u) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{u^2}{2}\right)$$

**Bandwidth (likelihood-optimal):**

$$h_k^* = \arg \max_h \sum_{i=1}^n \log \hat{p}_{-i}(\theta_k^{(i)} | h)$$

**Bandwidth (heuristic used in practice):**

$$h_k = c \cdot \text{std}\left(\{\theta_k^{(i)}\}_{i=1}^n\right), \quad c > 0$$